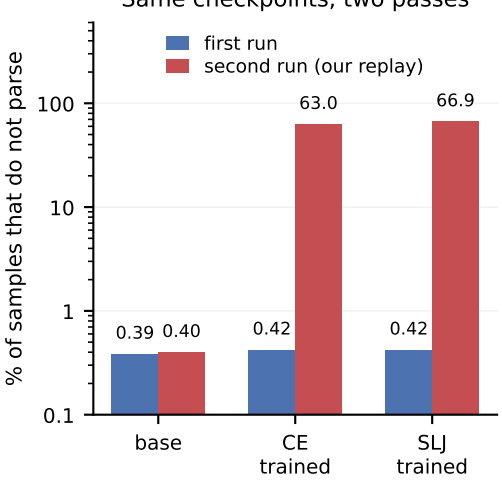
Same checkpoints, two passes
first run
second run (our replay)
% of samples that do not parse
100
10
1
0.1
0.39
0.40
0.42
63.0
0.42
66.9
base
CE trained
SLJ trained

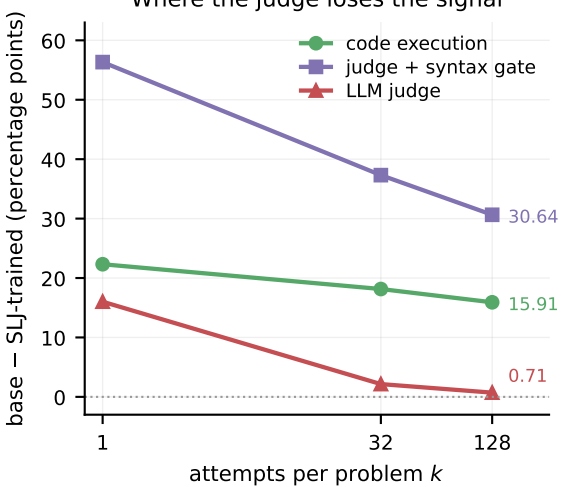
Where the judge loses the signal
code execution
judge + syntax gate
LLM judge
base − SLJ-trained (percentage points)
60
50
40
30
20
10
0
30.64
15.91
0.71
1
32
128
attempts per problem k